

MS SQL Server DBA

Praveen Kumar M
Mb: +91 986 613 0093 (Botim\WhatsApp)
+91 819 729 3434 (WhatsApp)

Mail: praveensqldba12@gmail.com

LinkedIn: <https://www.linkedin.com/in/praveenmadupu>

Github: <https://github.com/praveenmadupu>

Youtube: <https://www.youtube.com/PraveenMadupu>

Windows and SQL Server Failover Architecture

Windows and SQL Server Failover Architecture in Detail

Failover clustering in both **Windows Server** and **SQL Server** is a critical component for ensuring **high availability** (HA) and **disaster recovery** in mission-critical environments. The failover architecture involves setting up a system that allows services to remain operational if a hardware or software failure occurs on one node (server). This is achieved by configuring multiple servers (nodes) within a cluster, and if one node fails, another node takes over, minimizing downtime and service interruptions.

In this architecture, **Windows Server Failover Clustering** (WSFC) is the underlying technology that ensures high availability at the operating system level, while **SQL Server Failover Cluster Instances** (FCI) provide database-level high availability.

Let's explore both **Windows Failover Clustering (WSFC)** and **SQL Server Failover Cluster Instances (FCI)** architectures in detail:

1. Windows Server Failover Clustering (WSFC)

Windows Server Failover Clustering (WSFC) is the foundation for high availability in Windows environments. It allows multiple servers to work together as a cluster, providing redundancy and failover capabilities.

Components of Windows Server Failover Cluster (WSFC)

1. Cluster Nodes:

- A **node** in a failover cluster is simply a physical or virtual server that participates in the cluster.
- Nodes are typically dedicated to running critical workloads (like SQL Server, Exchange, etc.).
- A failover cluster can have multiple nodes, and one node can take over the responsibilities of another if it fails.
- Nodes in a WSFC are connected through a private network (for cluster communication).

2. Shared Storage:

- Shared storage is a critical component of a failover cluster and is typically implemented using a **SAN (Storage Area Network)** or **NAS (Network Attached Storage)**.
- All nodes in the cluster need access to the same shared disk. This shared storage holds the application data, such as SQL Server database files, so that they can be accessed by any node in the event of a failover.

3. Cluster Quorum:

- The **quorum** is a mechanism that ensures that the cluster can maintain a majority decision-making process in the event of a failure or partition (split-brain scenario).
- A quorum is required to avoid scenarios where two nodes may think they are the active node. There are several types of quorum configurations:
 - **Node Majority:** Requires an odd number of nodes and assumes that a majority of nodes are online.
 - **Disk Majority:** Requires the quorum disk to be online, and a majority of nodes to be online.
 - **Node and Disk Majority:** Combines both node and disk-based quorum.
 - **Cloud Witness/File Share Witness:** A witness server or cloud resource is used to act as a tie-breaker if there is an even number of nodes.

4. Cluster Network:

- The cluster network connects all the nodes in the cluster and is used for communication between nodes.
- The cluster uses both **public** (for client communication) and **private** (for internal cluster communication) networks.

5. Cluster Resource:

- Each node in the cluster manages resources, which could be virtual machines, databases, or other services. These resources are **grouped** and can fail over to another node if one node goes offline.

6. Failover Process:

- In case of a failure on one node, WSFC automatically detects the failure and triggers the failover process.
- The services or applications (like SQL Server) running on the failed node are moved to the next available node in the cluster.
- The **Failover Cluster Manager** is used to monitor and manage the failover cluster.

2. SQL Server Failover Cluster Instance (FCI)

A **SQL Server Failover Cluster Instance (FCI)** provides high availability for SQL Server by leveraging Windows Server Failover Clustering. FCI ensures that the SQL Server database engine and its associated data are protected from hardware or operating system failures.

Components of SQL Server Failover Cluster Instance (FCI)

1. SQL Server Failover Cluster Instance (FCI):

- An **FCI** is a single instance of SQL Server that is installed across multiple nodes in a failover cluster.
- All nodes in the cluster share access to the same **SQL Server databases** and can take over the SQL Server instance if another node fails.
- **FCI** allows the SQL Server instance to be accessed from the same network name regardless of which node in the cluster is currently active.

2. Cluster Shared Storage:

- Like WSFC, SQL Server FCI uses **shared storage** (such as a SAN or NAS) to store the SQL Server database files (data, log, and tempdb).
- The shared storage is accessible by all nodes in the cluster, allowing failover to occur seamlessly without data loss.
- If the active node fails, the SQL Server instance on another node can access the shared storage and take over without needing to resynchronize data.

3. Virtual Server Name:

- An FCI uses a **virtual server name** and **virtual IP address**. The virtual server name allows client applications to connect to SQL Server without needing to know which node is active.
- The virtual server name is part of the SQL Server instance, and when a failover occurs, the virtual server name moves to the new active node, allowing clients to reconnect to the SQL Server instance automatically.

4. SQL Server Instance Failover:

- If the active node hosting the SQL Server instance goes down, the failover process starts. The cluster manager automatically detects the failure and moves the SQL Server instance to another healthy node in the cluster.
- The cluster keeps track of which node holds the SQL Server instance, and this information is updated in the **SQL Server Error Log**.
- **Automatic Failover** ensures minimal downtime as SQL Server will automatically restart on the next available node.

5. SQL Server Services in an FCI:

- The SQL Server instance services (SQL Server Database Engine, SQL Server Agent, etc.) are installed on

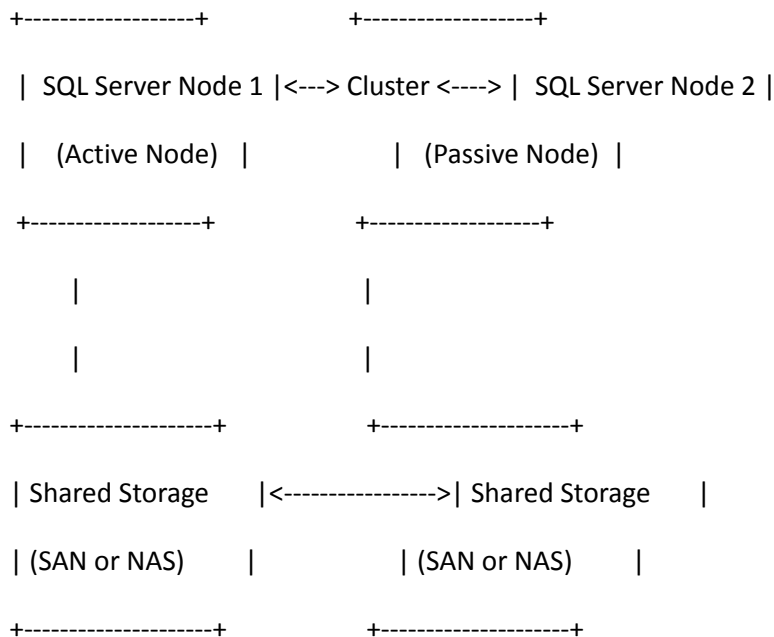
the **active node** and are moved to a **passive node** during failover.

- These services are configured as cluster resources, meaning they are highly available across all nodes in the cluster.

6. Transaction Log Shipping for Failover Clustering:

- For larger systems, **Transaction Log Shipping** can also be combined with failover clustering to ensure data is replicated to a secondary node.
- This is not a requirement for SQL Server FCI, but in large deployments, it ensures that in the event of a failover, the secondary server is fully up to date.

SQL Server Failover Clustering Architecture Overview



Failover Process in SQL Server FCI:

1. **Normal Operation:** SQL Server runs on one node (e.g., Node 1) and services requests from clients.
2. **Failure Detection:** If Node 1 experiences a failure (e.g., hardware, OS crash), the failover cluster detects the failure.
3. **Failover:** The **Failover Cluster Manager** triggers the failover process, and Node 2 assumes the role of the active node.
4. **SQL Server Restarts:** SQL Server services are started on Node 2, and the SQL Server instance becomes available to users with minimal downtime.
5. **Client Connection:** The client applications reconnect to the SQL Server instance using the **virtual server name** and virtual IP address, which are now pointing to Node 2.

Advantages of SQL Server Failover Clustering

1. High Availability:

- Automatic failover ensures minimal downtime if a node or SQL Server instance fails.
- With shared storage, the failover is seamless, and data remains consistent without needing to reinitialize the entire system.

2. Disaster Recovery:

- In case of hardware failure, SQL Server FCI ensures that the service continues with no data loss.
- Failover can occur automatically or be manually triggered by administrators.

3. **Centralized Management:**

- The entire cluster can be managed through **Failover Cluster Manager** and SQL Server Management Studio (SSMS), allowing centralized control over the entire system.

4. **Support for Multiple Nodes:**

- Failover clusters can support up to 16 nodes (depending on the SQL Server and Windows version), which offers scalability and high fault tolerance.

Disadvantages of SQL Server Failover Clustering

1. **Single Instance Limitation:**

- SQL Server FCI supports a single instance per cluster, which may limit scalability for large environments where multiple instances are required.

2. **Shared Storage Requirement:**

- FCI requires shared storage, which adds complexity and potential costs (i.e., SAN or NAS). Additionally, shared storage must be highly available, requiring additional redundancy.

3. **Complex Setup and Maintenance:**

- Setting up an FCI requires careful planning, and there can be a steep learning curve to understand failover clusters and ensure everything is properly configured and monitored.

4. **Hardware and Software Dependencies:**

- Failover clustering relies on specific hardware and software configurations (such as Windows Server versions, SANs, etc.), and compatibility must be ensured for a successful setup.

Summary:

Both **Windows Server Failover Clustering (WSFC)** and **SQL Server Failover Cluster Instances (FCI)** provide a robust framework for achieving **high availability** in SQL Server environments. While **WSFC** ensures the availability of the underlying operating system and services, **SQL Server FCI** specifically protects the SQL Server instance itself.

SQL Server FCI is an essential component for mission-critical applications that cannot afford significant downtime, offering automatic failover, data protection, and centralized management. However, setting up and maintaining a failover cluster involves careful planning, configuration, and hardware considerations. For large-scale environments, other solutions like **AlwaysOn Availability Groups** may be more appropriate, but FCI remains a reliable option for SQL Server high availability.